

# Report

## Assignment 2 - MySQL

Group: 35

Students: Albert Lesniewski, Thomas Tran, Md Anwarul Hasan

### Introduction

In this assignment we were tasked with designing a database schema for the GeoLife dataset, importing the data into the database, and doing any necessary data processing, and finally, running queries against the database in order to answer some of the questions that were specified as part of the assignment. We used the default proposed schema in the problem text. We focused on making our import code work in batches in order to achieve good import speed. As a group, we have worked mostly together on campus, in a pair-programming style, where we were programming and discussing the tasks at the same time. This way everyone could contribute to the programming tasks and be familiar with all parts of the code. For the questions that were supposed to be answered by querying the database, we split those between group members evenly and everyone did their part.

### Results

#### Part 1

For part 1 we were supposed to import the GeoLife dataset into mysql database. We used the default proposed schema, as it seemed fairly reasonable. Meaning we first defined queries for creating the tables (user, activity and track\_point). First we imported data for the user database, then the activities, and lastly the track points.

User table contains only two columns, the user id, and whether the user has specified the labels for their activities. The user id is a simple integer and the has\_labels column is determined by the python script by checking if a given user id is specified in the labeled\_id.txt file. Below you can see the top 10 entries in the user table after importing the data.

```

+-----+-----+
| id | has_labels |
+-----+-----+
| 000 | 0 |
| 001 | 0 |
| 002 | 0 |
| 003 | 0 |
| 004 | 0 |
| 005 | 0 |
| 006 | 0 |
| 007 | 0 |
| 008 | 0 |
| 009 | 0 |
+-----+-----+
10 rows in set (0.00 sec)

```

Activity table contains 5 columns, the activity id, which is just a auto incremented integer, user id which works as a foreign key constraint and references the id of a user in the user table, transportation mode which specifies what kind of transportation was used for a activity, this column can be NULL, in which case no transportation mode was specified, start date time and end date time are column which specify the start and end date and time for an activity. Filling the activity table with data was probably the most challenging, as it required quite some data processing. Figuring out the transportation mode was especially difficult. In our implementation we built a hash map (python dictionary) for each user, with activity id as key and the transportation mode as value. This way, when building the insert query, we would check the hashmap, and check if the activity id for the current user exists in the hash map. If it exists, we add the specified mode to the query. If not, the value is set to NULL. When building the insert query we iterate over all users. For each user we get that user's activities, and build the insert values for that user. In the end we concatenate all the insert values for all the users into one insert query and we execute it. This way we do only one insert query for all activities instead of executing the insert query hundreds or thousands of times. Below is the top 10 entries in the activity table.

```

+-----+-----+-----+-----+-----+
| id | user_id | transportation_mode | start_date_time | end_date_time |
+-----+-----+-----+-----+-----+
| 1 | 135 | NULL | 2009-01-03 01:21:34 | 2009-01-03 05:40:31 |
| 2 | 135 | NULL | 2009-01-02 04:31:27 | 2009-01-02 04:41:05 |
| 3 | 135 | NULL | 2009-01-27 03:00:04 | 2009-01-27 04:50:32 |
| 4 | 135 | NULL | 2009-01-10 01:19:47 | 2009-01-10 04:42:47 |
| 7 | 135 | NULL | 2009-01-14 12:17:57 | 2009-01-14 12:30:53 |
| 8 | 135 | NULL | 2009-01-12 01:41:22 | 2009-01-12 02:14:01 |
| 11 | 135 | NULL | 2008-12-24 14:42:07 | 2008-12-24 15:26:45 |
| 12 | 135 | NULL | 2008-12-28 10:36:05 | 2008-12-28 12:19:32 |
| 14 | 132 | NULL | 2010-02-15 10:56:35 | 2010-02-15 12:22:33 |
| 16 | 132 | NULL | 2010-04-30 23:38:01 | 2010-05-01 00:35:31 |
+-----+-----+-----+-----+-----+
10 rows in set (0.00 sec)

```

Track point table contains 7 columns. Track point id, which is auto incremented integer, activity id, which binds (references) a track point to a activity, latitude of the track point, longitude of the track point, altitude of the track point, date days is number of days since the UNIX epoch and lastly date time of the track point. As specified in the task, we have filtered out activities that contain more than 2500 track points. We have tried to insert all the track points as one insert query, like we did with activities, but we got a problem with exceeding the max packet limit. Thus we split the insert queries into batches of 1000 insert values at a time. Below you can see the top 10 entries in the track point table.

id	activity_id	lat	lon	altitude	date_days	date_time
4355470	1	39.974294	116.399741	492	39816.0566435185	2009-01-03 01:21:34
4355471	1	39.974292	116.399592	492	39816.0566550926	2009-01-03 01:21:35
4355472	1	39.974309	116.399523	492	39816.0566666667	2009-01-03 01:21:36
4355473	1	39.97432	116.399588	492	39816.0566898148	2009-01-03 01:21:38
4355474	1	39.974365	116.39973	491	39816.0567013889	2009-01-03 01:21:39
4355475	1	39.974391	116.399782	491	39816.0567361111	2009-01-03 01:21:42
4355476	1	39.974426	116.399735	491	39816.0567824074	2009-01-03 01:21:46
4355477	1	39.974458	116.3997	491	39816.0568402778	2009-01-03 01:21:51
4355478	1	39.974491	116.399732	490	39816.0568981481	2009-01-03 01:21:56
4355479	1	39.97453	116.399758	489	39816.0569560185	2009-01-03 01:22:01

10 rows in set (0.00 sec)

On a MacBook with a m1 chip, we managed to insert all the data in around 5 minutes.

## Part 2

Q1: In order to answer how many users, activities, and track points there are in the table we executed those queries:

```
SELECT COUNT(*) FROM user;
SELECT COUNT(*) FROM activity;
SELECT COUNT(*) FROM track_point;
```

The results where:

```
Number of users in table:
COUNT(*)
-----
      182
Number of activities in table:
COUNT(*)
-----
    16048
Number of track points in table:
COUNT(*)
-----
   9681756
```

**Q2:** In order to find the average number of activities per user, the number of activities and users where found. The average was then calculated by dividing the number of activities by the number of users.

```
def task2(self):
    query_user = "SELECT COUNT(*) FROM user"
    query_activity = "SELECT COUNT(*) FROM activity"

    self.cursor.execute(query_user)
    no_users = self.cursor.fetchall()

    self.cursor.execute(query_activity)
    no_activities = self.cursor.fetchall()

    average = no_activities[0][0] / no_users[0][0]

    print("The average number of activities per user is {} ".format(average))
```

```
The average number of activities per user is 88.17582417582418
```

**Q3:** To find the top 20 users with the highest number of activities, we run queries in the activity table. We counted the activity of each user and selected the top 20.

```
def task1(self):
    query1 = """SELECT user_id, count(user_id) as activity_count FROM activity
                group by user_id order by activity_count DESC LIMIT 20"""
    self.cursor.execute(query1)
    rows = self.cursor.fetchall()
    print(tabulate(rows, headers=self.cursor.column_names))
```

user_id	activity_count
128	2102
153	1793
025	715
163	704
062	691
144	563
041	399
085	364
004	346
140	345
167	320
068	280
017	265
003	261
014	236
126	215
030	210
112	208
011	201
039	198

Q4: In order to find out all users that have taken a taxi we have executed the following query:

```
SELECT user_id FROM activity WHERE transportation_mode='taxi' GROUP
BY user_id;
```

The results where:

```
Users that have taken a taxi:
user_id
-----
    010
    058
    062
    078
    080
    085
    098
    111
    128
    163
```

Q5: The number of activities per transportation mode was found by using the keyword **DISTINCT** and the count of a given transportation mode. In our activity table, those activities that are not labeled have a value of **NULL**.

```
def task5(self):
    query = """
    SELECT DISTINCT transportation_mode, COUNT(transportation_mode) AS count
    FROM activity
    WHERE transportation_mode IS NOT NULL
    GROUP BY transportation_mode
    """
    self.cursor.execute(query)
    rows = self.cursor.fetchall()
    print(tabulate(rows, headers=self.cursor.column_names))
```

transportation_mode	count
walk	480
bike	263
bus	199
subway	133
taxi	37
car	419
train	2
run	1
airplane	3
boat	1

Q6:

```
def task1(self):
    query1 = """SELECT year(start_date_time) as year, count(user_id) as
    total_no_activity FROM activity group by year
    order by total_no_activity desc limit 1"""
    self.cursor.execute(query1)
    rows = self.cursor.fetchall()
    print(tabulate(rows, headers=self.cursor.column_names))

    query2 = """SELECT year(start_date_time) as year, sum(hour(timediff(end_date_time,
    start_date_time))) as Total_Hours FROM activity
    group by year order by Total_Hours desc"""
    self.cursor.execute(query2)
    rows = self.cursor.fetchall()
    print(tabulate(rows, headers=self.cursor.column_names))
```

- a) To find the year with the most activity we searched the activity table. We selected each year & counted the number of activities each year. Thus we found the year with the most activity.

year	total_no_activity
2008	5895

- b) To get the answer whether this year has the most recorded hours we calculated total hours of each year and found that this is not the year that has the most recorded hours rather year 2009 has the most recorded hours.

c)

year	Total_Hours
2009	9165
2008	6921
2007	1915
2010	745
2011	688
2012	479
2000	0

Q7: In order to find the total distance walked by user 112 in 2008, we first had to execute a query that would give us all the track points for user 112 that were registered in 2008 and for which the activity was registered with walk as the transportation mode, the query was:

```
SELECT activity.id, lat, lon
```

```
FROM activity INNER JOIN track_point ON
activity.id=track_point.activity_id
WHERE user_id=112 AND YEAR(start_date_time)=2008 AND
transportation_mode='walk'
ORDER BY activity.id
```

Then, we wrote a simple python script that processed the resulting track point data from the query, and used the haversine distance to compute the distance between track points that are registered in latitude and longitude coordinates.

The result was:

```
Total distance walked by user 112 in 2008 is: 114.92625793252165km
```

Q8: To find the top 20 users who have gained the most altitude, the user table, activity table and a sub table of activity and track\_point containing the altitude difference for each activity were inner joined so that the program could take the sum of all the altitudes in every activity connected to each user. Because the altitude is given in feet and the solution requires meters, the altitude was multiplied by 0.3048.

	id	total_meters_gained_per_user
	128	350170
	153	268689
	144	140473
	041	130463
	062	121529
	163	96386
	004	83949.8
	025	83322.9
	085	73327.6
	003	57919.9
	140	54951.5
	039	44626.7
	030	43966.2
	167	41329.4
	002	37298.4
	034	36930.5
	084	35847.8
	126	35739.6
	042	33072
	037	32638.6

Q9: To find the users with invalid activity and also total number of invalid activity, we first searched for track points which have a gap of 5 minutes between them. We used LAG and TIMESTAMPDIFF in this regard. Then we combined these results with the activity table to find the users with invalid activity by comparing the id field of the activity table with activity\_id field of trackpoint table. After that we summed the invalid activity and grouped those by user\_id. The query and the outputs are provided below.



```
def task9(self):

    query = """SELECT q2.user_id, COUNT(q2.aid) FROM (
        SELECT q.user_id, q.aid FROM (
            SELECT activity.id aid, user_id, IFNULL(TIMESTAMPDIFF(minute, LAG(date_time) OVER (ORDER BY track_point.id ASC), date_time), 0) diff
                FROM activity INNER JOIN track_point ON activity_id=activity.id) as q
            INNER JOIN user ON user.id=q.user_id
        WHERE q.diff > 5
        GROUP BY q.aid) as q2
    GROUP BY q2.user_id
    """
```

User with invalid activity	Number of invalid activities		
135	6	040	17
132	3	078	61
104	96	047	9
103	33	065	53
168	50	091	79
157	11	096	62
150	19	062	467
159	6	054	2
166	4	053	9
161	12	098	5
102	23	038	64
105	9	007	31
133	4	000	125
134	53	009	35
158	10	036	41
167	223	031	3
151	1	052	49
169	20	099	13
024	39	055	17
023	14	063	9
015	50	097	28
012	55	090	6
079	14	064	13
046	26	030	165
041	290	008	19
048	1	037	115
077	3	001	52
083	24	039	164
084	129	006	20
070	6	174	65
013	69	180	3
014	174	173	5
022	68	145	5
025	454	142	101
071	40	129	5
085	269	116	1
082	50	111	30
076	11	118	5
		127	6

144	364	011	122
172	14	016	28
181	15	029	29
175	4	081	19
121	4	075	8
119	28	072	2
126	158	086	5
110	23	044	44
128	1387	088	39
117	5	043	28
153	1156	017	199
154	22	028	43
162	10	010	59
165	8	026	20
131	11	019	55
136	13	021	7
109	3	003	225
100	6	004	277
107	1	032	12
138	14	035	24
164	7	095	22
163	456	061	16
155	32	066	6
152	2	092	121
106	3	059	5
139	16	050	8
101	61	057	19
108	7	068	214
130	11	034	117
089	49	033	4
042	77	005	58
045	8	002	127
087	5	056	18
073	60	069	6
074	49	051	47
080	7	093	15
020	86	067	60
027	2	058	14
018	35	060	1

  

094	19
112	146
115	122
123	3
124	4
170	2
141	3
146	7
179	33
125	32
122	6
114	15
113	12
147	55
140	206
176	7
171	4

Q10: In order to find out what users have registered an activity in the Forbidden City in Beijing we had to execute the following query:

```
SELECT user_id
FROM activity INNER JOIN track_point ON
activity.id=track_point.activity_id
WHERE lat BETWEEN 39.915 AND 39.917 AND lon BETWEEN 116.396
AND 116.398
GROUP BY user_id
```

The result was:

```
Users that have registered an activity in the Forbidden City (threshold of 0.001 in latitude and longitude):
user_id
-----
131
018
019
004
```

**Q11:** To find the most used transportation mode for each user (which has their activities labeled), the program first constructs a view of a table with the user and count for each transportation mode the user has. This view is then later used in another query. In this query, the user id and the max transportation mode is chosen so that only one transportation mode is chosen if a user has the same number of activities tagged with the labels. user\_tmode (which is the view) is then joined with another subquery (called T) which contains the number of counts of the transportation mode that occurs the most for a user. By doing this, the most used transportation mode for each user is found.

id	most_used_transportation_mode
010	taxi
020	bike
021	walk
052	bus
056	bike
058	walk
060	walk
062	walk
064	bike
065	bike
067	walk
069	bike
073	walk
075	walk
076	car
078	walk
080	taxi
081	bike
082	walk
084	walk
085	walk
086	car
087	walk
089	car
091	walk
092	walk
097	bike
098	taxi
101	car
102	bike
107	walk
108	walk
111	taxi
112	walk
115	car
117	walk
125	bike
126	bike
128	car
136	walk
138	bike
139	bike
144	walk
153	walk
161	walk
163	bike
167	bike
175	bus

```
def task11(self):
    query = """
    CREATE OR REPLACE VIEW user_tmode AS
    SELECT DISTINCT user.id, transportation_mode, COUNT(transportation_mode) AS count_tm
    FROM user
    INNER JOIN activity ON user.id = activity.user_id
    WHERE transportation_mode IS NOT NULL
    GROUP BY user.id, transportation_mode
    """

    self.cursor.execute(query)

    query = """
    SELECT user_tmode.id, MAX(user_tmode.transportation_mode) AS most_used_transportation_mode
    FROM user_tmode
    JOIN (SELECT id, MAX(count_tm) AS max_count_tm FROM user_tmode GROUP BY id) T
    ON user_tmode.count_tm = T.max_count_tm AND user_tmode.id = T.id
    GROUP BY user_tmode.id
    ORDER BY user_tmode.id ASC
    """

    self.cursor.execute(query)

    rows = self.cursor.fetchall()

    print(tabulate(rows, headers=self.cursor.column_names))
```

---

## Discussion

In part 1 of this assignment, what surprised us the most was how few of the labels actually matched with an activity. But the reason as to why this happened might be due to the method we used to match them. We looked at the start time and end time for a label and tried to match it up with the start time and end time of an activity by looking at its first and last track point. However, there might have been a few seconds or even minutes differences which we didn't take into account. Removing activities with more than 2500 track points might have affected this result as well, because some of these activities might've had labels.

In part 2, something that we didn't expect was the amount of invalid activities from query 9. By comparing the number of invalid activities of an user with an user's total number of activities, we saw that in general, almost half of an user's activities were invalid.

From this assignment, by writing the data in batches to the database, we saw the practical results of sequential writes vs random writes. By writing in batches, the data is written in clumps of blocks, which makes it faster than if it were to write each data by its own across a storage medium. This was clearly shown, as we got a huge performance boost when writing the data in batches.

## Summary

In this assignment, we managed to clean up the data from the dataset and insert it into a database. By using a combination of python and SQL queries, we also retrieved the desirable results as described in the query tasks.